

BIBLIOMETRIC ANOMALY DETECTION

Tier 2b :: Citation
Rings & Collusion-
Level Misconduct

At a glance: Bibliometric anomaly detection identifies journals, authors, or institutions whose citation and publication patterns deviate significantly from field-level norms — flagging potential citation rings, coercive citation, editorial manipulation, or paper mill activity. It operates on metadata available from academic databases and does not require access to the underlying manuscripts.

1. Overview

Every researcher, journal, and institution leaves a statistical footprint in the academic record: patterns of who cites whom, how often, and in what context. Under ordinary conditions these patterns reflect the genuine intellectual structure of a field — seminal papers are widely cited, recent publications build on recent work, and citation exchange between groups is roughly proportional to intellectual proximity. Bibliometric anomaly detection asks whether observed patterns deviate from what those normal conditions would predict.

The approach has roots in the sociology of science (Merton, de Solla Price) and became practically powerful as large-scale bibliometric databases — Web of Science, Scopus, OpenAlex, Semantic Scholar — made it possible to compute network statistics across millions of papers. What was once a laborious manual exercise can now be automated across entire journal catalogues.

Unlike tools such as GRIM or statcheck, which examine individual papers for impossible statistics, bibliometric anomaly detection operates at the journal, author, or institutional level. It does not identify fabricated data directly — it identifies structural patterns that are inconsistent with legitimate scholarly activity and that warrant further investigation. A flagged journal or author cluster is a starting point for inquiry, not a finding of misconduct.

2. Methodological Overview

Several distinct bibliometric signals are used in combination. No single metric is sufficient on its own — the diagnostic value comes from the convergence of multiple anomalies, or from the severity of a single anomaly relative to a well-calibrated field baseline.

Self-citation rate

Self-citation — authors citing their own prior work — is normal and legitimate up to a point. The anomaly threshold varies by field: in highly specialized niches, rates above 20–25% are common; in broad fields, rates above 10–15% are suspicious. Extreme self-citation (> 50%) can artificially inflate h-indices and impact metrics and is flagged by several databases. The key metric is whether the self-citation rate significantly exceeds the distribution of comparable authors in the same subfield and career stage.

Co-citation and reciprocal citation networks

Citation network analysis maps directed edges between papers. Anomalies of interest include:

- Dense reciprocal citation clusters: groups of authors who disproportionately cite each other relative to their citation of the broader field. Network modularity and clustering coefficients identify these sub-communities.
- Implausible citation velocity: a paper accumulates citations far faster than its intellectual predecessors in the same line of work, or than papers of comparable type from the same period.
- Citation concentration: a journal's impact factor is driven by a small number of papers that are cited extremely heavily relative to the rest of the journal's output — consistent with editorial self-promotion or cite-stacking.

Journal-level metrics

Several journal-level statistics can be computed from database metadata and compared to field distributions:

Self-citation ratio = citations_from_same_journal / total_citations_received

Ratios above 0.3–0.4 are flagged by ESCI and Scopus review committees as potentially manipulated. Clarivate has suppressed impact factors for journals exceeding this threshold.

Co-citation density = intra-cluster citations / total citations (normalized by cluster size)

Compares observed citation density within an author or journal cluster to the expected density under a null model of random citation with the same degree distribution.

Additional journal-level signals include: unusually rapid peer review turnaround times (consistent with fake or superficial review); implausibly high acceptance rates; concentration of authorship in a narrow geographic or institutional cluster; and discordance between stated and actual editorial board membership.

Author-level metrics

For individual authors, anomaly detection examines:

- h-index inflation: an h-index that substantially exceeds what would be expected given the author's publication record, career length, and subfield, suggesting inflated citation through a network.
- Productivity outliers: paper output rates that significantly exceed field norms and are not explained by large-team or consortium authorship. Paper mills generate authorship credits at implausible rates.
- Authorship network anomalies: an author appears across many papers in fields with no apparent intellectual connection to their stated expertise — a signature of purchased or gifted authorship.

3. Example Application: Journal Impact Factor Manipulation

Clarivate Analytics has suppressed the Journal Citation Reports impact factor for multiple journals following bibliometric audit. A well-documented case involved a cluster of journals in the cell biology and chemistry space that engaged in coordinated cite-stacking: editors solicited or required authors to cite a specific set of papers from the same journal as a condition of acceptance.

The bibliometric signature was distinctive: the journals' self-citation ratios were between 0.35 and 0.60 — far above the field baseline of 0.08–0.12 — and the cited papers were concentrated in a small number of issues rather than distributed across the journal's history. Network analysis showed that the citing papers were often methodologically unrelated to the cited papers, confirming the citations were instrumental rather than substantive.

Clarivate suppressed the impact factors of the implicated journals, removing them from the Journal Citation Reports for two years. Several were ultimately delisted from the Science Citation Index entirely. The case illustrated both the power of bibliometric anomaly detection — the pattern was identifiable without reading any individual paper — and its limitation: some journals engaged in coordinated self-citation for years before the statistical signal crossed the detection threshold.

A second class of examples involves paper mills. Richardson & Amaral (PNAS 2025) applied network analysis to PLOS ONE editorial and authorship data, finding that 0.25% of editors were responsible for 30.2% of retracted articles. The bibliometric signal — concentration of editorial throughput in a tiny fraction of the editor network, combined with anomalous authorship clustering in those submissions — allowed identification of the implicated editors before full forensic review of the manuscripts themselves.

4. Requirements & Applicability

Bibliometric anomaly detection requires access to structured citation metadata. The quality and completeness of that metadata is the primary constraint on what can be detected.

Citation database access	
	Web of Science, Scopus, or OpenAlex (free) provide the citation metadata needed. OpenAlex is fully open-access via API and covers > 250 million works. Semantic Scholar provides similar coverage. Some analyses require institutional database licences.

Defined comparison baseline	Anomalies are only detectable relative to a baseline. This requires either a reference set of comparable journals or authors in the same subfield, or a statistical null model of expected citation behavior. The choice of baseline is a significant analyst decision.
Network analysis tools	Citation network construction and community detection require tools such as VOSviewer, Gephi, or network analysis packages in R (igraph, tidygraph) or Python (NetworkX). Bibliometrix (R) provides a complete pipeline from database export to network statistics.
Database export format	Web of Science and Scopus export structured reference data in formats that bibliometric tools can ingest directly. OpenAlex provides a REST API. Raw DOI or title lists are insufficient — full reference metadata is required for network construction.
Field-level norms	Interpretation requires knowing what normal self-citation, citation velocity, and network density look like in the target field. These vary substantially — high in mathematics and physics, lower in biomedical research. Anomalies must be evaluated against field-specific distributions, not universal thresholds.
Time window definition	Citation accumulation is time-dependent. Analyses must specify a publication year window and a citation observation period. Comparing citation rates across different time windows produces spurious anomalies.
No manuscript access required	The method operates entirely on citation and authorship metadata. Full text of the papers is not needed, which makes it scalable to large literature sets.

5. Limitations

- Bibliometric anomalies are consistent with misconduct but not diagnostic of it. A dense citation cluster may reflect a legitimate intellectual community. High self-citation may reflect a researcher at the frontier of a narrow specialization. Anomaly detection generates hypotheses for further investigation, not findings.
- Field norms vary enormously and are poorly documented for many subfields. An anomaly in one field may be standard practice in another. Applying generic thresholds (e.g., 'self-citation > 30% is suspicious') across fields generates both false positives and false negatives.
- Database coverage is uneven. OpenAlex and Semantic Scholar are rapidly improving but still undercount citations in regional-language journals, conference proceedings (especially outside computer science), and pre-2000 literature. Anomalies in poorly covered fields may be artefacts of missing data.
- Legitimate citation incentives mimic the anomaly signature. Impact factor pressure encourages authors to cite journals they want to publish in, generating organic cite-stacking without editorial coercion. Distinguishing intentional manipulation from rational response to incentives is difficult.
- Network methods depend heavily on the chosen community detection algorithm and resolution parameters. Different algorithms applied to the same citation network can produce very different community boundaries, changing which clusters appear anomalous.
- The method is retrospective and slow. Citation patterns take years to become statistically distinguishable. Rapidly expanding paper mill operations may not trigger anomaly flags until hundreds of fraudulent papers have already been published and cited.

6. How to Apply

Step 1 — Define scope and baseline

Specify the target entity (journal, author, institution) and the field. Export a reference set of comparable journals or authors from the same Web of Science or Scopus category. This reference set forms the empirical baseline for anomaly thresholds.

Step 2 — Export citation metadata

From Web of Science or Scopus, export full reference data (not just abstract records) for the target entity and the reference set. In OpenAlex, use the Works API with filters on source, author, or institution. Retain all metadata fields: DOI, references, authors, affiliations, publication date.

Step 3 — Compute self-citation and journal-level metrics

Using Bibliometrix (R) or the bibliometrix web app (bibliometrix.org), import the exported data and compute: self-citation ratio, h-index, citation distribution, and co-citation matrix. Compare values to the reference set distribution. Flag entities in the top or bottom 5th percentile on multiple metrics simultaneously.

Step 4 — Build and analyze the citation network

Construct a directed citation network where nodes are papers or authors and edges are citations. Use igraph (R) or NetworkX (Python) to compute: modularity, clustering coefficient, betweenness centrality. Apply community detection (Louvain or Leiden algorithm). Visualise with VOSviewer or Gephi. Flag communities with internal edge density substantially higher than expected under a configuration null model.

Step 5 — Cross-reference with retraction and alert databases

Check flagged journals and authors against the Retraction Watch Database, Clarivate's suppressed journal list, and Beall's List of predatory publishers. Convergence of bibliometric anomalies with known misconduct indicators substantially increases the probability of genuine problems.

7. Related Tools

Citation Network Analysis	Specifically maps the directed graph of citation relationships to identify clusters and community structure. Bibliometric anomaly detection draws on citation network methods but adds statistical comparison to field baselines.
Editor Network Analysis	Applies network methods to editorial submission routing rather than citation patterns. Richardson & Amaral (PNAS 2025) used this to identify editors enabling fake peer review at scale — a structural complement to citation-level anomaly detection.
Problematic Paper Screener	Operates at the manuscript level, flagging tortured phrases and image duplication. Convergence between a Problematic Paper Screener flag and a bibliometric anomaly on the same journal provides much stronger evidence of systematic misconduct than either alone.
Hijacked Journal Checker	Identifies journals whose ISSNs have been cloned by scam publishers. A prerequisite check before bibliometric analysis — if a journal is hijacked, its citation metadata may be split across the legitimate and fraudulent versions, distorting metrics.
Retraction Watch Database	Provides the ground truth for calibrating bibliometric anomaly thresholds: journals with high retraction rates can be used as positive controls when building detection models.

Key Citations

- Richardson, R. A. K., Hong, S. S., Byrne, J. A., Stoeger, T., & Amaral, L. A. N. (2025). The entities enabling scientific fraud at scale are large, resilient, and growing rapidly. *PNAS*, 122(32), e2420092122. <https://doi.org/10.1073/pnas.2420092122>
- Smaldino, P. E. & McElreath, R. (2016). The natural selection of bad science. *Royal Society Open Science*, 3(9), 160384. <https://doi.org/10.1098/rsos.160384>
- Biagioli, M. & Lippman, A. (Eds.). (2020). *Gaming the metrics: Misconduct and manipulation in academic research*. MIT Press.
- Aria, M. & Cuccurullo, C. (2017). bibliometrix: An R-tool for comprehensive science mapping analysis. *Journal of Informetrics*, 11(4), 959–975. <https://doi.org/10.1016/j.joi.2017.08.007>
- Fister, I., Jr., Fister, I., & Perc, M. (2016). Toward the discovery of citation cartels in citation networks. *Frontiers in Physics*, 4, 49. <https://doi.org/10.3389/fphy.2016.00049>
- Clarivate Analytics. (2023). *Journal Citation Reports: Suppressed journals policy*. <https://clarivate.com/products/scientific-and-academic-research/research-analytics-evaluation-and-management-solutions/journal-citation-reports>

Harzing, A. W. & Alakangas, S. (2016). Google Scholar, Scopus and the Web of Science: A longitudinal and cross-disciplinary comparison. *Scientometrics*, 106(2), 787–804. <https://doi.org/10.1007/s11192-015-1798-9>
